

For the 2 C values I attempted, I did the specified $\sqrt{2}$ and chose 100 as my own...

For $C = \sqrt{2}$, the AI's decisions were more balanced and seemed more human-like. This value is theorized to have the best amount of balance between exploration and exploitation. So, the results make sense, given this fact.

When I ran the code with $C = 100$, the AI more often chose to move a piece it hadn't moved yet. I think I also noticed that it tended to repeat moves it had conducted if it was moving the same piece (barring jumps vs normal moves). This makes sense since a higher C value favors exploration vs exploitation, making it more likely to attempt new moves.

Additionally, I feel like the AI was harder to play against when $C = \sqrt{2}$ because it was more unpredictable. Whereas I felt like I could gauge what the AI would do when $C = 100$.